

Do Disclosed Institutional Trades Predict Returns?

Evidence from 48,879 NSE bulk and block deals in Indian large caps

Praveen Kumar · nse-smartmoney project · July 2026

Data: NSE bulk/block-deal disclosures 2024 to 2026, delivery statistics, daily OHLCV

Abstract

A popular idea among retail investors is that you can profit by following institutional “smart money.” In India this is unusually testable, because exchanges publish every large trade under the name of the client behind it. I checked the idea against 48,879 disclosed deals from January 2024 to July 2026, matched to two years of daily prices for the NIFTY 50.

At first the folklore looked right. One foreign desk’s repeated purchases of a large-cap stock were followed by a 1.98% average 10-day return, a 92% win rate, and a one-sided p-value of 0.0003. That is exactly the kind of result that gets screenshotted and shared.

It did not survive scrutiny. Once I adjusted for the market’s own move over the same days, the excess return fell to roughly zero. The twelve “independent” events turned out to be a single three-week buying campaign, run at the same time across several index names, so they carried no stock-specific information. After a false-discovery-rate correction, nothing was left. Pooled machine-learning models trained on delivery, volume and deal-flow features scored below a coin flip out of sample, with a ROC-AUC of 0.46 to 0.48. My conclusion is that in Indian large caps, the simple version of smart-money tracking is mostly an illusion. I also set out where a fairer test should look next: mid-cap stocks, campaign-level event definitions, and longer flow histories.

1. Introduction

Indonesian traders call it bandarmology, the art of reading large operators from broker-level flow. India gives you a sharper lens. SEBI rules require exchanges to disclose, every day and by name, two kinds of trade. A bulk deal is any single client trading more than 0.5% of a company’s equity. A block deal is a large negotiated window trade. Where an Indonesian analyst has to guess who sits behind broker code “BK,” the Indian analyst is simply told. The name might be Goldman Sachs Bank Europe SE, or SBI Funds Management, or a private individual.

That transparency invites a specific question. When a named institution buys the same stock again and again, does the stock then go up? This paper reports what I found when I put that question through a disciplined statistical pipeline. The headline is a negative result, and I think it is more useful to a practitioner than another success story would have been.

2. Data

All of the data come from public NSE sources. I loaded them into a SQLite warehouse using an open-source pipeline I wrote, called nse-smartmoney, which handles the ETL, feature engineering, event studies and dashboard.

Dataset	Coverage	Size	Source
Bulk-deal disclosures	Jan 2024 to Jul 2026	46,177 deals	NSE report archive (manual CSV)

			export)
Block-deal disclosures	Jan 2025 to Jun 2026	2,702 deals	NSE report archive (manual CSV export)
Daily OHLCV, NIFTY 50	Jul 2024 to Jul 2026	25,000 sym-days	Yahoo Finance
Security-wise delivery %	Jan 2026 to Jul 2026	272,440 sym-days	NSE bhavcopy
FII/DII aggregate flows	Jul 2026 (2 days)	4 rows	NSE (accumulating daily)

The 48,879 deals span 2,443 different symbols. Only 844 of them, or 1.7%, touch a NIFTY 50 constituent. That is itself a finding. The 0.5% ownership threshold makes bulk deals rare in mega-caps, so most disclosed smart-money activity actually happens in mid- and small-cap names that sit outside this study's price universe.

I classify each deal participant into a behavioural profile using keyword rules over the disclosed name.

Profile	Deals	Gross value (Rs crore)
Other / Retail-HNI	27,401	767,178
Proprietary / HFT desk	9,820	291,296
FII / FPI	8,195	519,449
DII - AIF / PMS / NBFC	1,837	51,760
DII - Mutual Fund	738	112,975
HNI / Family Office	732	14,096
DII - Insurance	156	9,395

3. Methodology

- Participant event study. For every stock-and-participant pair with at least five distinct net-buy days, I compute forward 5- and 10-day returns from each event close, then test whether the mean is greater than zero with a one-sided t-test at the 5% level.
- Robustness cascade. Three follow-up checks. First, market adjustment: subtract the equal-weight universe forward return over the same windows. Second, a clustering control: look at how the event dates are spaced, and collapse consecutive-day events into a single campaign when the gap is under seven calendar days. Third, multiple testing: a Benjamini-Hochberg adjustment across every pair tested.
- Delivery-conviction study. A delivery-percentage z-score over a 60-day window, confirmed by volume, defines an accumulation day. I then run a Welch t-test of forward returns on spike days versus ordinary days.
- Pooled predictability. An OLS regression of 5-day forward returns on delivery, volume, momentum and deal-flow features, with Newey-West errors. On top of that, a logistic regression and a random forest with a chronological 70/30 split, scored by out-of-sample ROC-AUC.

Every feature is knowable at the close of the signal day. Forward returns are used only as labels. No look-ahead is allowed anywhere in the pipeline.

4. Results

4.1 The naive result looks strong

Five stock-and-participant pairs cleared the minimum-event threshold in the NIFTY 50 universe. Three of them passed the naive significance test. All three involved the same participant, the offshore-derivative (ODI) desk of Goldman Sachs Bank Europe SE.

Stock	Participant	Events	Mean 10-day fwd	Win rate	p (one-sided)
ETERNAL	Goldman Sachs Bank Europe SE - ODI	12	+1.98%	92%	0.0003
KOTAKBANK	Goldman Sachs Bank Europe SE - ODI	6	+2.32%	100%	0.0239
SHRIRAMFIN	Goldman Sachs Bank Europe SE - ODI	7	+3.75%	71%	0.0421
BHARTIARTL	Goldman Sachs Bank Europe SE - ODI	6	-0.25%	50%	0.5835
M&M	PI Opportunities AIF V LLP	6	-0.66%	33%	0.6213

Read at face value, this is the story that retail smart-money content is made of. A 92% win rate, a p-value of 0.0003, and a famous institution attached to it. It is also, as the next section shows, almost entirely an illusion.

4.2 The audit

The first problem is event clustering. The twelve ETERNAL “events” fall on thirteen trading days between 29 August and 15 September 2025. That is one continuous buying streak plus a single later purchase. Because the 10-day forward windows overlap, the observations are not independent, and the effective sample is closer to two campaigns than twelve events. KOTAKBANK and SHRIRAMFIN are worse still. Their event dates are identical to each other, running from 29 August to 12 September 2025. All three “signals” are the same basket-buying program printed across index names at once. That is the fingerprint of an index or ODI flow, not of stock selection. Across the whole sample this one desk traded 274 disclosed deals in 121 different symbols.

The second problem is the market. Once I subtract the equal-weight universe return over the same windows, the picture changes.

Stock	Raw 10-day mean	Market-adjusted mean	Adjusted win rate	Adjusted p
ETERNAL	+1.16%	-0.10%	54%	0.542
KOTAKBANK	+2.41%	+1.19%	71%	0.123
SHRIRAMFIN	+3.75%	+2.53%	71%	0.031

The flagship ETERNAL result, $p = 0.0003$ in the naive test, is statistically zero once the market’s own September 2025 rally is removed. The stock rose with everything else while the desk happened to be printing bulk deals.

The third problem is multiple testing. Under a Benjamini-Hochberg correction across the five tested pairs, only ETERNAL survives at q below 0.05, with an adjusted p of 0.0014. But ETERNAL is exactly the pair the other two checks invalidate. SHRIRAMFIN, the one pair that survives market adjustment at $p = 0.031$, then fails the FDR correction with an adjusted p of 0.070, and it rests on a single campaign of seven overlapping windows anyway.

So the verdict is blunt. No stock-and-participant pair in the NIFTY 50 universe survives all three checks.

4.3 Pooled models find no generic edge

If smart-money footprints carried systematic information, a pooled model should pick it up. It does not. On a panel of 4,450 symbol-days from February to July 2026, where every feature is observable, the results are flat.

Model / test	Key statistic	Result
OLS, HAC errors (5-day fwd returns)	R-squared	0.003
momentum (5-day back return)	coef +0.047, t = 2.46	p = 0.014 (weak positive)
smart-money deal net quantity	t = -3.30	small negative sign
delivery z, volume z, accum. score	-	all insignificant
Logistic regression (OOS, n = 1,335)	ROC-AUC	0.461
Random forest (OOS, n = 1,335)	ROC-AUC	0.477

Both classifiers score below a coin flip out of sample. One detail is worth flagging. The deal-flow coefficient is slightly negative, which means that on average, large-cap days with disclosed institutional buying were followed by marginally weaker returns. That fits a world where a disclosure marks a liquidity event, an exit or a rebalance, rather than a conviction entry.

4.4 Delivery spikes, if anything, lean the wrong way

Across 79 volume-confirmed delivery-spike days in the first half of 2026, the mean 5-day forward return was -0.72%, against +0.03% on ordinary days, with a Welch t of -1.90 and p = 0.061. Individual symbols are noisy and two-sided on very small samples. JSW Steel comes out positive at p = 0.046, Tata Consumer negative at p below 0.001. The popular heuristic that high delivery means accumulation, and accumulation means buy, gets no support here. If anything the point estimate leans the other way.

5. Discussion

- Large caps hide footprints by design. A bulk deal needs 0.5% of a company's equity in one day. In a company worth ten lakh crore rupees, that is an enormous position. Only 1.7% of disclosed deals touch NIFTY 50 names, and the ones that do are dominated by basket, index and ODI flows that carry no stock-specific view.
- The unit of analysis matters a lot. Counting each day of a multi-day buying program as a separate event manufactures significance out of nothing. Campaigns, not days, are the honest unit. In this sample most pairs amount to one or two campaigns.
- Market adjustment is not optional. Every naive positive in this study rode a rising market. Excess return over the index, not raw return, is the least you should demand as evidence.
- A negative result is still informative. This does not prove smart-money tracking is worthless in India. It proves the large-cap version of the folklore fails, and it points to exactly where a real test has to go.

6. Limitations

- The price history covers NIFTY 50 only. That leaves 98.3% of the disclosed deals untested.
- Only five pairs reached the minimum-event threshold, so the participant-level test is underpowered in this universe by construction.

- Delivery data span six months, a single market regime. The FII/DII flow history is only two days old and grows daily.
- Bulk-deal history for the first half of 2026 was still downloading when I wrote this. The 2024 and 2025 files are complete.
- Participant classification is a keyword heuristic, and ODI desks trade for undisclosed end clients.

7. Future work

- Extend the price universe to the NIFTY Midcap 150 and Smallcap 250, where the 48,000-deal history is actually concentrated and event counts per participant are far higher.
- Re-run the audit with campaigns as the event unit and market-adjusted returns as the default metric.
- Keep collecting FII/DII flow and delivery history through the pipeline's daily scheduled job, then revisit the flow-causality (Granger) tests once there is six months or more of it.
- Track specific high-conviction domestic participants, such as mutual funds and AIFs, whose mandates imply real stock selection rather than basket execution.

8. Conclusion

India's disclosure regime is unusually transparent, so the smart-money question deserves a real answer rather than folklore. The answer I got is sobering. In the NIFTY 50, every naive institutional-footprint signal I found dissolved under three standard checks. Pooled models found nothing a coin flip would miss, and the delivery heuristic pointed the wrong way. The upside is the infrastructure. The pipeline, warehouse, validation layer and dashboard I built for this audit now collect data every day, which is what will let me ask the sharper version of the question where it belongs, in the mid-cap names where disclosed institutions actually leave tracks.

Disclaimer: this paper is educational research, not investment advice. Bulk and block-deal disclosures identify the executing counterparty, not intent or the end client, and nothing here alleges wrongdoing by any named institution. Analysis code and data pipeline: github.com/Pravkum23/nse-smartmoney.